

Algorithmic Fairness on Graphs: Methods and Trends

Jian Kang

University of Illinois at Urbana-Champaign
jiank2@illinois.edu

Hanghang Tong

University of Illinois at Urbana-Champaign
htong@illinois.edu

ABSTRACT

Graph is a ubiquitous type of data that appears in many real-world applications, including social network analysis, recommendations and financial security. Important as it is, decades of research have developed plentiful computational models to mine graphs. Despite its prosperity, concerns with respect to the potential algorithmic discrimination have been grown recently. Algorithmic fairness on graphs, which aims to mitigate bias introduced or amplified during the graph mining process, is an attractive yet challenging research topic. The first challenge corresponds to the theoretical challenge, where the non-IID nature of graph data may not only invalidate the basic assumption behind many existing studies in fair machine learning, but also introduce new fairness definition(s) based on the inter-correlation between nodes rather than the existing fairness definition(s) in fair machine learning. The second challenge regarding its algorithmic aspect aims to understand how to balance the trade-off between model accuracy and fairness. This tutorial aims to (1) comprehensively review the state-of-the-art techniques to enforce algorithmic fairness on graphs and (2) enlighten the open challenges and future directions. We believe this tutorial could benefit researchers and practitioners from the areas of data mining, artificial intelligence and social science.

CCS CONCEPTS

• Information systems → Data mining.

KEYWORDS

Graph mining, algorithmic fairness, bias

ACM Reference Format:

Jian Kang and Hanghang Tong. 2022. Algorithmic Fairness on Graphs: Methods and Trends. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD '22)*, August 14–18, 2022, Washington, DC, USA. ACM, New York, NY, USA, 2 pages. <https://doi.org/10.1145/3534678.3542599>

1 AUDIENCE PARTICIPATION

The tutorial is designed for researchers and practitioners in data mining, artificial intelligence, social science and other interdisciplinary areas. The prerequisites include basic knowledge on probability, linear algebra, data mining and machine learning. No prior knowledge on algorithmic fairness or any specific algorithm is required. This tutorial is designed for 40% novice, 30% intermediate,

30% expert, in order to achieve a good balance between the introductory and advanced materials. During the tutorial, we will alter between lecture and discussion with the audience to encourage audience participation. We will also post the materials of this tutorial (e.g., tutorial description, presentation slides, pre-recorded videos) for post-tutorial dissemination.

2 PRESENTER BIOGRAPHY

Jian Kang is currently a Ph.D. student in the Department of Computer Science at the University of Illinois at Urbana-Champaign. He received his M.CS. degree in Computer Science from the University of Virginia in 2016 and B.Eng. degree in Telecommunication Engineering from Beijing University of Posts and Telecommunications in 2014. His current research interests lie in large-scale data mining and machine learning, especially on graphs, with a focus on their algorithmic fairness. His research works on related topics have been published at several major conferences and journals in data mining and artificial intelligence. He has also served as a reviewer and a program committee member in top-tier data mining and artificial intelligence venues and journals (e.g., NeurIPS, ICML, ICLR, CIKM, WSDM, JMLR, TKDE, etc). For more information, please refer to his personal website at <http://jiank2.web.illinois.edu/>.

Hanghang Tong is currently an associate professor at Department of Computer Science at University of Illinois at Urbana-Champaign. Before that he was an associate professor at School of Computing, Informatics, and Decision Systems Engineering (CIDSE), Arizona State University. He received his M.Sc. and Ph.D. degrees from Carnegie Mellon University in 2008 and 2009, both in machine learning. His research interest is in large-scale data mining for graphs and multimedia. He has received several awards, including SDM/IBM Early Career Data Mining Research award (2018), NSF CAREER award (2017), ICDM 10-Year Highest Impact Paper award (2015), and four best paper awards (TUP'14, CIKM'12, SDM'08, ICDM'06). He has published over 100 refereed articles. He is the Editor-in-Chief of SIGKDD Explorations (ACM), an action editor of Data Mining and Knowledge Discovery (Springer), and an associate editor of Knowledge and Information Systems (Springer); and has served as a program committee member in data mining, database and artificial intelligence venues (e.g., SIGKDD, SIGMOD, AAAI, WWW, CIKM). He is a Fellow of IEEE. For more information, please refer to his personal website at <http://tonghanghang.org/>.

3 OUTLINE OF TUTORIAL

• Introduction

We will discuss with the background knowledge for the topic, then briefly review existing problem definitions and settings, along with the key challenges in this topic. We will also introduce problems in relation to algorithmic fairness on graphs, including auditing [6], adversarial attack [17], privacy preservation [16] and uncertainty quantification [15].

- Background and motivations

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

KDD '22, August 14–18, 2022, Washington, DC, USA

© 2022 Copyright held by the owner/author(s).

ACM ISBN 978-1-4503-9385-0/22/08.

<https://doi.org/10.1145/3534678.3542599>

- Problem definitions and settings
- Key challenges
- Related problems

• Part I: Group Fairness on Graphs

In this part, we will discuss how to enforce fairness among nodes of different demographic groups, i.e., group fairness, for classic graph mining (e.g., ranking, clustering) and for advanced graph mining (e.g., graph neural networks).

- Classic graph mining [12, 13]
- Advanced graph mining [2, 3, 8, 10]

• Part II: Individual Fairness on Graphs

In this part, we will introduce how to operationalize individual fairness on graphs.

- Optimization-based methods [5]
- Ranking-based methods [4]

• Part III: Other Fairness Notions on Graphs

In this part, we will study how to enforce three fairness notions on graphs as follows, other than the extensively studied group fairness and individual fairness.

- Rawlsian fairness [7]
- Degree-related fairness [11]
- Counterfactual fairness [1, 9]

• Part IV: Beyond Fairness on Graphs

In this part, we will discuss problems in relation to fairness, as well as their corresponding state-of-the-art techniques.

- Interpretability [14]
- Robustness [17]
- Accountability [6]

• Part V: Future Trends

In this part, we elucidate open challenges and future directions from the following perspectives.

- Benchmark datasets and evaluation metrics
- Interpretable fairness on graphs
- Fairness on dynamic graphs

4 RELATED TUTORIALS

Here, we discuss the most relevant tutorials in other conferences, as well as the similarities/differences compared with ours.

• Fair Graph Mining

- **Presenters:** Jian Kang, Hanghang Tong
- **Conference:** CIKM, Nov. 1 - 5, 2021, Virtual Conference
- **Connection:** This is an earlier version of this tutorial on algorithmic fairness on graph mining.
- **Difference:** There will be some overlaps between the related tutorial and this tutorial (e.g., preliminaries and foundation of fair graph algorithms). However, in this tutorial, we will identify new challenge(s) (e.g., edge-level fairness), introduce recent advances that were not covered before (e.g., [3, 7–9, 12]) and provide new insights on the future directions (e.g., benchmark datasets and evaluation metrics for algorithmic fairness on graphs, relationship between fairness and other important aspects of graph algorithms such as interpretability, robustness, etc.)

• Fairness in Networks

- **Presenters:** Sorelle Friedler, Carlos Scheidegger, Suresh Venkatasubramanian, Aaron Clauset
- **Conference:** KDD, Aug. 14 - 18, 2021, Virtual Conference

- **Connection:** Both tutorials aim to introduce fairness on graphs.
- **Difference:** The related tutorial focuses on fairness from the social welfare perspective for information access and influence maximization, whereas our tutorial present fairness on graphs in a much broader scope. Our tutorial differs the related tutorial in two aspects: (1) we review various fairness notions including group fairness, individual fairness, degree-related fairness and counterfactual fairness, in addition to the fairness from the social welfare aspect (e.g., the Rawlsian difference principle); (2) we include a wider coverage of research works in graph mining, including classic graph mining (e.g., ranking and clustering) and advanced graph mining (e.g., embedding and graph neural networks).

5 POTENTIAL SOCIETAL IMPACTS

This tutorial has several potential positive impacts to the society:

- (1) we hope this tutorial could attract research attention to promote fairness on graphs, which is less popular than fairness for IID data;
- (2) we hope this tutorial raises new challenges that are not addressed in the existing works;
- (3) we hope this tutorial could also benefit related research topic in identifying new problems and discovering its relationship to fairness on graphs.

ACKNOWLEDGEMENT

This work is supported by NSF (1947135, 2134079 and 1939725).

REFERENCES

- [1] Chirag Agarwal, Himabindu Lakkaraju, and Marinka Zitnik. 2021. Towards a Unified Framework for Fair and Stable Graph Representation Learning. *arXiv preprint arXiv:2102.13186* (2021).
- [2] Avishek Bose and William Hamilton. 2019. Compositional Fairness Constraints for Graph Embeddings. In *ICML*.
- [3] Anyan Dai and Suhang Wang. 2021. Say No to the Discrimination: Learning Fair Graph Neural Networks with Limited Sensitive Attribute Information. In *WSDM*.
- [4] Yushun Dong, Jian Kang, Hanghang Tong, and Jundong Li. 2021. Individual Fairness for Graph Neural Networks: A Ranking based Approach. In *KDD*.
- [5] Jian Kang, Jingrui He, Ross Maciejewski, and Hanghang Tong. 2020. InFoRM: Individual Fairness on Graph Mining. In *KDD*. 379–389.
- [6] Jian Kang and Hanghang Tong. 2019. N2N: Network Derivative Mining. In *CIKM*. 861–870.
- [7] Jian Kang, Yan Zhu, Yinglong Xia, Jiebo Luo, Hanghang Tong, and Xin Wang. 2022. RawlsGCN: Towards Rawlsian Difference Principle on Graph Convolutional Network. In *WWW*.
- [8] Peizhao Li, Yifei Wang, Han Zhao, Pengyu Hong, and Hongfu Liu. 2020. On Dyadic Fairness: Exploring and Mitigating Bias in Graph Connections. In *ICLR*.
- [9] Jing Ma, Ruocheng Guo, Mengting Wan, Longqi Yang, Aidong Zhang, and Jundong Li. 2022. Learning Fair Node Representations with Graph Counterfactual Fairness. *WSDM*.
- [10] Tahleel A Rahman, Bartłomiej Surma, Michael Backes, and Yang Zhang. 2019. Fairwalk: Towards Fair Graph Embedding. In *IJCAL*. 3289–3295.
- [11] Xianfeng Tang, Huaxiu Yao, Yiwei Sun, Yiqi Wang, Jiliang Tang, Charu Aggarwal, Prasenjit Mitra, and Suhang Wang. 2020. Investigating and Mitigating Degree-Related Biases in Graph Convolutional Networks. In *CIKM*. 1435–1444.
- [12] Sotiris Tsioutsoulis, Evaggelia Pitoura, Konstantinos Semertzidis, and Panayiotis Tsaparas. 2022. Link Recommendations for PageRank Fairness. In *WWW*.
- [13] Sotiris Tsioutsoulis, Evaggelia Pitoura, Panayiotis Tsaparas, Ilias Klefakis, and Nikos Mamoulis. 2021. Fairness-Aware PageRank. In *WWW*. 3815–3826.
- [14] Hao Yuan, Haiyang Yu, Jie Wang, Kang Li, and Shuiwang Ji. 2021. On explainability of graph neural networks via subgraph explorations. In *ICML*. 12241–12252.
- [15] Xujiang Zhao, Feng Chen, Shu Hu, and Jin-Hee Cho. 2020. Uncertainty Aware Semi-Supervised Learning on Graph Data. In *NeurIPS*, Vol. 33.
- [16] Jun Zhou, Chaochao Chen, Longfei Zheng, Xiaolin Zheng, Bingzhe Wu, Ziqi Liu, and Li Wang. 2020. Privacy-Preserving Graph Neural Network for Node Classification. *arXiv preprint arXiv:2005.11903* (2020).
- [17] Daniel Zügner, Amir Akbarnejad, and Stephan Günnemann. 2018. Adversarial Attacks on Neural Networks for Graph Data. In *KDD*. 2847–2856.